Capstone Proposal

REVIEW | CODE REVIEW | HISTORY

Meets Specifications

Dear student

Great job on your proposal! I think that you're meeting the specifications and it's clear that you have a pretty good idea of what you want to do. Your suggestions are all feasible and I think you're on the right track.

About the capstone project:
While the code and implementation are both important, keep in mind that the capstone report is the most important element of your final project. This report simulates a formal submission to a journal for peer-review. Publishing your results is a key credential in machine learning and we want you to be ready for this!

You should have very little trouble quickly passing the project review if you initially follow the report template. Keep in mind that reviewers will be looking to see that you not only fully document how you implemented your project, but why you made the choices you made. This type of critical thinking is crucial to publishing in a peer-reviewed journal. Based on your proposal, I don't think you'll have much trouble with this, but I wanted to mention it up front.

I think you're definitely on solid ground and you've picked an interesting topic for your project. I'm looking forward to seeing the final result!

Cheers!

Project Proposal

✓ **Student clearly describes the problem that is to be solved. The problem is well defined and has at least one relevant potential solution. Additionally, the problem is quantifiable, measurable, and replicable.**

> The problem can be quantified in the following terms: number of current/established customer clusters (customer segmentation unsupervised problem) and probability of being a new customer to the company (supervised problem).

I think that this is a great summary of the problem and how it is structured. Nice work!

✓ **Student briefly details background information of the domain from which the project is proposed. Historical information relevant to the project should be included. It should be clear how or why a problem in the domain can or should be solved. Related academic research should be appropriately cited. A discussion of the student's personal motivation for investigating a particular problem in the domain is encouraged but not required.**

Great job giving the reader an introduction to the problem domain!

Suggested:

- If you include a link to your data source in this section, you can directly lift it into the `Project Overview` section of your capstone report (to save a bit of time).

✓ **Student clearly describes a solution to the problem. The solution is applicable to the project domain and appropriate for the dataset(s) or input(s) given. Additionally, the solution is quantifiable, measurable, and replicable.**

You've given a great high-level summary of your solution. Nice work!

✓ **Student proposes at least one evaluation metric that can be used to quantify the performance of both the benchmark model and the solution model presented. The evaluation metric(s) proposed are appropriate given the context of the data, the problem statement, and the intended solution.**

> To be more precise, the ranking/scoring is based on AUC, the curve being the ROC curve.

AUC is a good metric for this dataset since the dataset classes are imbalanced. Be sure to avoid using accuracy for this project since this can introduce bias into your models:
https://towardsdatascience.com/accuracy-paradox-897a69e2dd9b

✓ **A benchmark model is provided that relates to the domain, problem statement, and intended solution. Ideally, the student's benchmark model provides context for existing methods or known information in the domain and problem given, which can then be objectively compared to the student's solution. The benchmark model is clearly defined and measurable.**

> Thus, our benchmark model will be a standard Logistic Regression model with outcomes 1 = new customer, 0 = not a new customer.

Nice choice! This is a common default implementation that won't be too hard to beat.

✓ **The dataset(s) and/or input(s) to be used in the project are thoroughly described. Information such as how the dataset or input is (was) obtained, and the characteristics of the dataset or input, should be included. It should be clear how the dataset(s) or input(s) will be used in the project and whether their use is appropriate given the context of the problem.**

This is a solid summary of the scope and nature of the dataset. Well done!

Suggested:

- Don't forget to note how the dataset classes are balanced in your project report. This is also an important consideration for why you should avoid using accuracy as a metric with this dataset.

✓ **Student summarizes a theoretical workflow for approaching a solution given the problem. A discussion is made as to what strategies may be employed, what analysis of the data might be required, or which algorithms will be considered. The workflow and discussion provided align with the qualities of the project. Small visualizations, pseudocode, or diagrams are encouraged but not required.**

Suggested:

- The CatBoost and LightGBM models could be good complementary ensemble supervised learning approaches to try here.
- Since you're creating multiple supervised learning models, you could try combining them all together into a custom ensemble model:

http://blog.kaggle.com/2016/12/27/a-kagglers-guide-to-model-stacking-in-practice/
https://www.kaggle.com/arthurtok/introduction-to-ensembling-stacking-in-python

✓ **The proposal follows a well-organized structure and would be readily understood by its intended audience. Each section is written in a clear, concise and specific manner. Few grammatical and spelling mistakes are present. All resources used and referenced are properly cited.**

The template format is followed and the proposal is well written.

DOWNLOAD PROJECT

RETURN TO PATH

Rate this review